linear_regression

January 22, 2018

0.1 Linear regression workbook

This workbook will walk you through a linear regression example. It will provide familiarity with Jupyter Notebook and Python. Please print (to pdf) a completed version of this workbook for submission with HW #1.

ECE 239AS, Winter Quarter 2018, Prof. J.C. Kao, TAs C. Zhang and T. Xing

```
In [134]: import numpy as np
          import matplotlib.pyplot as plt

          #allows matlab plots to be generated in line
          %matplotlib inline
```

0.1.1 Data generation

For any example, we first have to generate some appropriate data to use. The following cell generates data according to the model: $y = x - 2x^2 + x^3 + \epsilon$

```
In [135]: np.random.seed(0)   # Sets the random seed.
          num_train = 200      # Number of training data points

          # Generate the training data
          x = np.random.uniform(low=0, high=1, size=(num_train,))
          y = x - 2*x**2 + x**3 + np.random.normal(loc=0, scale=0.03, size=(num_train,))
          f = plt.figure()
          ax = f.gca()
          ax.plot(x, y, '.')
          ax.set_xlabel('$x$')
          ax.set_ylabel('$y$')
```

```
Out[135]: <matplotlib.text.Text at 0x1126bf850>
```

0.1.2 QUESTIONS:

Write your answers in the markdown cell below this one:

(1) What is the generating distribution of x?

(2) What is the distribution of the additive noise ϵ?

0.1.3 ANSWERS:

(1) The generating distribution of x is Uniform

(2) The distribution of the additive noise is Normal

0.1.4 Fitting data to the model (5 points)

Here, we'll do linear regression to fit the parameters of a model $y = ax + b$.

```
In [136]: # xhat = (x, 1)
          xhat = np.vstack((x, np.ones_like(x)))

          # ==================== #
          # START YOUR CODE HERE #
          # ==================== #
          # GOAL: create a variable theta; theta is a numpy array whose elements are [a, b]
```

2

```
theta = np.zeros(2) # please modify this line
theta = np.linalg.inv((xhat).dot(xhat.T)).dot(xhat.dot(y))
# =================== #
# END YOUR CODE HERE #
# =================== #
```

In [137]:
```
# Plot the data and your model fit.
f = plt.figure()
ax = f.gca()
ax.plot(x, y, '.')
ax.set_xlabel('$x$')
ax.set_ylabel('$y$')

# Plot the regression line
xs = np.linspace(min(x), max(x),50)
xs = np.vstack((xs, np.ones_like(xs)))
plt.plot(xs[0,:], theta.dot(xs))
```

Out[137]: [<matplotlib.lines.Line2D at 0x112866fd0>]

0.1.5   QUESTIONS

(1) Does the linear model under- or overfit the data?

(2) How to change the model to improve the fitting?

0.1.6 ANSWERS

(1) The model is underfitting.

(2) We can improve the fitting by using an equation of higher order or degree i.e., to increase the dimensions of theta and X correspondingly.

0.1.7 Fitting data to the model (10 points)

Here, we'll now do regression to polynomial models of orders 1 to 5. Note, the order 1 model is the linear model you prior fit.

```
In [138]: N = 5
          xhats = []
          thetas = []

          # ===================== #
          # START YOUR CODE HERE #
          # ===================== #

          # GOAL: create a variable thetas.
          # thetas is a list, where theta[i] are the model parameters for the polynomial fit o
          #    i.e., thetas[0] is equivalent to theta above.
          #    i.e., thetas[1] should be a length 3 np.array with the coefficients of the x^2,
          #    ... etc.

          for i in np.arange(N):
              if i == 0:
                  thetas.append(theta)
                  xhats.append(xhat)
              else:
                  xhat = np.vstack((x**(i+1), xhat))
                  xhats.append(xhat)
                  thetas.append(np.linalg.inv((xhats[i]).dot(xhats[i].T)).dot(xhats[i].dot(y))

          # ================== #
          # END YOUR CODE HERE #
          # ================== #

In [139]: # Plot the data
          f = plt.figure()
          ax = f.gca()
          ax.plot(x, y, '.')
          ax.set_xlabel('$x$')
          ax.set_ylabel('$y$')

          # Plot the regression lines
```

```
plot_xs = []
for i in np.arange(N):
    if i == 0:
        plot_x = np.vstack((np.linspace(min(x), max(x),50), np.ones(50)))
    else:
        plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))
    plot_xs.append(plot_x)

for i in np.arange(N):
    ax.plot(plot_xs[i][-2,:], thetas[i].dot(plot_xs[i]))

labels = ['data']
[labels.append('n={}'.format(i+1)) for i in np.arange(N)]
bbox_to_anchor=(1.3, 1)
lgd = ax.legend(labels, bbox_to_anchor=bbox_to_anchor)
```


0.1.8 Calculating the training error (10 points)

Here, we'll now calculate the training error of polynomial models of orders 1 to 5.

```
In [141]: training_errors = []
          # ==================== #
          # START YOUR CODE HERE #
          # ==================== #
          # GOAL: create a variable training_errors, a list of 5 elements,
          # where training_errors[i] are the training loss for the polynomial fit of order i+1
          for i in np.arange(N):
```

```
        training_errors.append((1/float(len(y))) * np.linalg.norm(y - thetas[i].dot(xhats

        pass

        # ================== #
        # END YOUR CODE HERE #
        # ================== #

        print ('Training errors are: \n', training_errors)
```

('Training errors are: \n', [0.0034496094622164849, 0.0023371908575540567, 0.0020210892856459088

0.1.9 QUESTIONS

(1) What polynomial has the best training error?

(2) Why is this expected?

0.1.10 ANSWERS

(1) The polynomial of degree 5 has the best training error (i.e., the least training error.)

(2) This was expected because a polynomial of higher degree covers all the points in the training
data thus decreasing the training error to its least value.

0.1.11 Generating new samples and testing error (5 points)

Here, we'll now generate new samples and calculate testing error of polynomial models of orders
1 to 5.

```
In [144]: x = np.random.uniform(low=1, high=2, size=(num_train,))
          y = x - 2*x**2 + x**3 + np.random.normal(loc=0, scale=0.03, size=(num_train,))
          f = plt.figure()
          ax = f.gca()
          ax.plot(x, y, '.')
          ax.set_xlabel('$x$')
          ax.set_ylabel('$y$')

Out[144]: <matplotlib.text.Text at 0x112a28e50>
```

```
In [145]: xhats = []
          for i in np.arange(N):
              if i == 0:
                  xhat = np.vstack((x, np.ones_like(x)))
                  plot_x = np.vstack((np.linspace(min(x), max(x),50), np.ones(50)))
              else:
                  xhat = np.vstack((x**(i+1), xhat))
                  plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))

              xhats.append(xhat)

In [146]: # Plot the data
          f = plt.figure()
          ax = f.gca()
          ax.plot(x, y, '.')
          ax.set_xlabel('$x$')
          ax.set_ylabel('$y$')

          # Plot the regression lines
          plot_xs = []
          for i in np.arange(N):
              if i == 0:
                  plot_x = np.vstack((np.linspace(min(x), max(x),50), np.ones(50)))
              else:
                  plot_x = np.vstack((plot_x[-2]**(i+1), plot_x))
```

7

```
            plot_xs.append(plot_x)

        for i in np.arange(N):
            ax.plot(plot_xs[i][-2,:], thetas[i].dot(plot_xs[i]))

        labels = ['data']
        [labels.append('n={}'.format(i+1)) for i in np.arange(N)]
        bbox_to_anchor=(1.3, 1)
        lgd = ax.legend(labels, bbox_to_anchor=bbox_to_anchor)
```

```
In [147]: testing_errors = []

          # ===================== #
          # START YOUR CODE HERE #
          # ===================== #

          # GOAL: create a variable testing_errors, a list of 5 elements,
          # where testing_errors[i] are the testing loss for the polynomial fit of order i+1.
          for i in np.arange(N):

              testing_errors.append((1/float(len(y))) * np.linalg.norm(y - thetas[i].dot(xhats

          pass

          # =================== #
          # END YOUR CODE HERE #
```

```
# ================= #
```

```
print ('Testing errors are: \n', testing_errors)
```

('Testing errors are: \n', [0.063585238792311649, 0.10324532058417471, 0.012501394138930177, 0

0.1.12 QUESTIONS

(1) What polynomial has the best testing error?

(2) Why polynomial models of orders 5 does not generalize well?

0.1.13 ANSWERS

(1) Polynomial of degree 4 has the best Testing Error

(2) The polynomial of degree 5 is over fitting the training data. Although the training error is the least, the 5th degree polynomial does not generalize well as it tries to cover maximum number points in the training data.