

```

1 import numpy as np
2
3 class Softmax(object):
4
5     def __init__(self, dims=[10, 3073]):
6         self.init_weights(dims=dims)
7
8     def init_weights(self, dims):
9         """
10        Initializes the weight matrix of the Softmax classifier.
11        Note that it has shape (C, D) where C is the number of
12        classes and D is the feature size.
13        """
14        self.W = np.random.normal(size=dims) * 0.0001
15
16    def loss(self, X, y):
17        """
18        Calculates the softmax loss.
19
20        Inputs have dimension D, there are C classes, and we operate on minibatches
21        of N examples.
22
23        Inputs:
24        - X: A numpy array of shape (N, D) containing a minibatch of data.
25        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
26            that X[i] has label c, where 0 <= c < C.
27
28        Returns a tuple of:
29        - loss as single float
30        """
31
32        # Initialize the loss to zero.
33        loss = 0.0
34        num_samples = X.shape[0]
35        # ===== #
36        # YOUR CODE HERE:
37        # Calculate the normalized softmax loss. Store it as the variable loss.
38        # (That is, calculate the sum of the losses of all the training
39        # set margins, and then normalize the loss by the number of
40        # training examples.)
41        # ===== #
42        aY = X.dot(self.W.T)
43        for i in range(num_samples):
44
45            logSum = np.log(np.sum(np.exp(aY[i])))
46            classProb = aY[i, y[i]]
47            loss += -classProb + logSum
48
49        # ===== #
50        # END YOUR CODE HERE
51        # ===== #
52        loss = loss / num_samples
53        return loss
54
55    def loss_and_grad(self, X, y):
56        """
57        Same as self.loss(X, y), except that it also returns the gradient.
58
59        Output: grad -- a matrix of the same dimensions as W containing
60            the gradient of the loss with respect to W.
61        """
62
63        # Initialize the loss and gradient to zero.
64        loss = 0.0
65        grad = np.zeros_like(self.W)
66        num_samples = X.shape[0]
67        numClasses = self.W.shape[0]
68        # ===== #
69        # YOUR CODE HERE:
70        # Calculate the softmax loss and the gradient. Store the gradient
71        # as the variable grad.
72        # ===== #
73        aY = X.dot(self.W.T)
74        for i in range(num_samples):
75            logSum = np.log(np.sum(np.exp(aY[i])))
76            classProb = aY[i, y[i]]
77            loss += -classProb + logSum
78
79            for j in range(numClasses):

```

```

1/31/2018 /Users/vijayravi/Documents/UCLA/Coursework/2018Winter/neuralNetworks/homeworks/homework2/code/nndl/softmax.py
81         grad[j] += ((np.exp(X[i].dot(self.W.T)[j])/float(np.sum(np.exp(X[i].dot(self.W.T)))) - (y[i] == j)) * X[i]
82
83     loss = loss/num_samples
84     grad = grad/num_samples
85     # ===== #
86     # END YOUR CODE HERE
87     # ===== #
88
89     return loss, grad
90
91 def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
92     """
93     sample a few random elements and only return numerical
94     in these dimensions.
95     """
96
97     for i in np.arange(num_checks):
98         ix = tuple([np.random.randint(m) for m in self.W.shape])
99
100         oldval = self.W[ix]
101         self.W[ix] = oldval + h # increment by h
102         fxph = self.loss(X, y)
103         self.W[ix] = oldval - h # decrement by h
104         fxmh = self.loss(X,y) # evaluate f(x - h)
105         self.W[ix] = oldval # reset
106
107         grad_numerical = (fxph - fxmh) / (2 * h)
108         grad_analytic = your_grad[ix]
109         rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))
110         print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))
111
112 def fast_loss_and_grad(self, X, y):
113     """
114     A vectorized implementation of loss_and_grad. It shares the same
115     inputs and outputs as loss_and_grad.
116     """
117     loss = 0.0
118     grad = np.zeros(self.W.shape) # initialize the gradient as zero
119     # ===== #
120     # YOUR CODE HERE:
121     # Calculate the softmax loss and gradient WITHOUT any for loops.
122     # ===== #
123     aY = X.dot(np.transpose(self.W))
124     aExp = np.exp(aY)
125
126     aCol = np.sum(aExp, axis = 1)
127
128     logSum = np.log(aCol)
129     sum1 = -aY[np.arange(X.shape[0]), y]
130     loss = np.sum(sum1+logSum)/X.shape[0]
131
132     aNorm = aExp / np.transpose(np.matrix(aCol))
133
134
135     aNorm[np.arange(X.shape[0]),y] -= 1
136     grad = aNorm.T.dot(X)
137
138     grad = grad/X.shape[0]
139
140
141     # ===== #
142     # END YOUR CODE HERE
143     # ===== #
144
145     return loss, grad
146
147 def train(self, X, y, learning_rate=1e-3, num_iters=100,
148         batch_size=200, verbose=False):
149     """
150     Train this linear classifier using stochastic gradient descent.
151
152     Inputs:
153     - X: A numpy array of shape (N, D) containing training data; there are N
154         training samples each of dimension D.
155     - y: A numpy array of shape (N,) containing training labels; y[i] = c
156         means that X[i] has label 0 <= c < C for C classes.
157     - learning_rate: (float) learning rate for optimization.
158     - num_iters: (integer) number of steps to take when optimizing
159     - batch_size: (integer) number of training examples to use at each step.
160     - verbose: (boolean) If true, print progress during optimization.

```

```

161
162 Outputs:
163 A list containing the value of the loss function at each training iteration.
164 """
165 num_train, dim = X.shape
166 num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes
167
168 self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W
169
170 # Run stochastic gradient descent to optimize W
171 loss_history = []
172
173 for it in np.arange(num_iters):
174
175     X_batch = None
176     y_batch = None
177     indic = np.random.choice(num_train, batch_size)
178     X_batch = X[indic,:]
179     y_batch = y[indic]
180
181     # ===== #
182     # YOUR CODE HERE:
183     # Sample batch_size elements from the training data for use in
184     # gradient descent. After sampling,
185     # - X_batch should have shape: (dim, batch_size)
186     # - y_batch should have shape: (batch_size,)
187     # The indices should be randomly generated to reduce correlations
188     # in the dataset. Use np.random.choice. It's okay to sample with
189     # replacement.
190     # ===== #
191
192     # ===== #
193     # END YOUR CODE HERE
194     # ===== #
195
196     # evaluate loss and gradient
197
198     loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
199     loss_history.append(loss)
200
201     # ===== #
202     # YOUR CODE HERE:
203     # Update the parameters, self.W, with a gradient step
204     # ===== #
205
206     self.W -= learning_rate*grad
207
208     # ===== #
209     # END YOUR CODE HERE
210     # ===== #
211
212     if verbose and it % 100 == 0:
213         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
214
215 return loss_history
216
217 def predict(self, X):
218     """
219     Inputs:
220     - X: N x D array of training data. Each row is a D-dimensional point.
221
222     Returns:
223     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
224       array of length N, and each element is an integer giving the predicted
225       class.
226     """
227     y_pred = np.zeros(X.shape[0])
228     # ===== #
229     # YOUR CODE HERE:
230     # Predict the labels given the training data.
231     # ===== #
232     prod = np.dot(X, np.transpose(self.W))
233
234     y_pred = np.argmax(prod, axis = 1)
235     # ===== #
236     # END YOUR CODE HERE
237     # ===== #
238     return y_pred

```