

```

1 import numpy as np
2 import pdb
3
4 """
5 This code was based off of code from cs231n at Stanford University, and modified for ece239as at UCLA.
6 """
7 class SVM(object):
8
9     def __init__(self, dims=[10, 3073]):
10         self.init_weights(dims=dims)
11
12     def init_weights(self, dims):
13         """
14         Initializes the weight matrix of the SVM. Note that it has shape (C, D)
15         where C is the number of classes and D is the feature size.
16         """
17         self.W = np.random.normal(size=dims)
18
19     def loss(self, X, y):
20         """
21         Calculates the SVM loss.
22
23         Inputs have dimension D, there are C classes, and we operate on minibatches
24         of N examples.
25
26         Inputs:
27         - X: A numpy array of shape (N, D) containing a minibatch of data.
28         - y: A numpy array of shape (N,) containing training labels; y[i] = c means
29             that X[i] has label c, where 0 <= c < C.
30
31         Returns a tuple of:
32         - loss as single float
33         """
34
35         # compute the loss and the gradient
36         num_classes = self.W.shape[0]
37         num_train = X.shape[0]
38         loss = 0.0
39         a_mat = np.dot(X, np.transpose(self.W))
40         for i in np.arange(num_train):
41             # ===== #
42             # YOUR CODE HERE:
43             # Calculate the normalized SVM loss, and store it as 'loss'.
44             # (That is, calculate the sum of the losses of all the training
45             # set margins, and then normalize the loss by the number of
46             # training examples.)
47             # ===== #
48
49             for j in range(num_classes):
50                 if(j != y[i]):
51                     ajx = a_mat[i,j]
52                     ayx = a_mat[i, y[i]]
53                     loss += np.maximum(0, 1+ajx-ayx)
54             # ===== #
55             # END YOUR CODE HERE
56             # ===== #
57         loss = loss/num_train
58         return loss
59
60     def loss_and_grad(self, X, y):
61         """
62         Same as self.loss(X, y), except that it also returns the gradient.
63
64         Output: grad -- a matrix of the same dimensions as W containing
65             the gradient of the loss with respect to W.
66         """
67
68         # compute the loss and the gradient
69         num_classes = self.W.shape[0]
70         num_train = X.shape[0]
71         loss = 0.0
72         grad = np.zeros_like(self.W)

```

```

73 a_mat = np.dot(X, np.transpose(self.W))
74 for i in np.arange(num_train):
75     # ===== #
76     # YOUR CODE HERE:
77     # Calculate the SVM loss and the gradient. Store the gradient in
78     # the variable grad.
79     # ===== #
80     for j in range(num_classes):
81         if(j != y[i]):
82             ajx = a_mat[i,j]
83             ayx = a_mat[i, y[i]]
84             zj = 1+ajx-ayx
85             loss += np.maximum(0, zj)
86             grad[j] += (zj > 0) * X[i]
87             grad[y[i]] -= (zj > 0) * X[i]
88
89     # ===== #
90     # END YOUR CODE HERE
91     # ===== #
92
93     loss /= num_train
94     grad /= num_train
95
96     return loss, grad
97
98 def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
99     """
100     sample a few random elements and only return numerical
101     in these dimensions.
102     """
103
104     for i in np.arange(num_checks):
105         ix = tuple([np.random.randint(m) for m in self.W.shape])
106
107         oldval = self.W[ix]
108         self.W[ix] = oldval + h # increment by h
109         fxph = self.loss(X, y)
110         self.W[ix] = oldval - h # decrement by h
111         fxmh = self.loss(X,y) # evaluate f(x - h)
112         self.W[ix] = oldval # reset
113
114         grad_numerical = (fxph - fxmh) / (2 * h)
115         grad_analytic = your_grad[ix]
116         rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) + abs(grad_analytic))
117         print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical, grad_analytic, rel_error))
118
119 def fast_loss_and_grad(self, X, y):
120     """
121     A vectorized implementation of loss_and_grad. It shares the same
122     inputs and outputs as loss_and_grad.
123     """
124     loss = 0.0
125     grad = np.zeros(self.W.shape) # initialize the gradient as zero
126
127     # ===== #
128     # YOUR CODE HERE:
129     # Calculate the SVM loss WITHOUT any for loops.
130     # ===== #
131     a_mat = np.dot(X, np.transpose(self.W))
132     ayx = a_mat[np.arange(X.shape[0]), y]
133
134     zj = 1+ a_mat - np.matrix(ayx).T
135     zj[np.arange(X.shape[0]), y] -= 1
136
137     loss += np.maximum(0, zj)
138
139     loss = np.sum(loss)/X.shape[0]
140     # ===== #
141     # END YOUR CODE HERE
142     # ===== #
143
144
145

```

```

146 # ===== #
147 # YOUR CODE HERE:
148 # Calculate the SVM grad WITHOUT any for loops.
149 # ===== #
150 ind_mat = np.zeros_like(zj)
151 ind_mat[zj>0] = 1
152 sum1 = np.sum(ind_mat, axis = 1)
153 ind_mat[np.arange(X.shape[0]), y] = -sum1.T
154 grad = np.matmul(ind_mat.T, X)
155 grad = grad/X.shape[0]
156 # ===== #
157 # END YOUR CODE HERE
158 # ===== #
159
160 return loss, grad
161
162 def train(self, X, y, learning_rate=1e-3, num_iters=100,
163         batch_size=200, verbose=False):
164     """
165     Train this linear classifier using stochastic gradient descent.
166
167     Inputs:
168     - X: A numpy array of shape (N, D) containing training data; there are N
169         training samples each of dimension D.
170     - y: A numpy array of shape (N,) containing training labels; y[i] = c
171         means that X[i] has label 0 ≤ c < C for C classes.
172     - learning_rate: (float) learning rate for optimization.
173     - num_iters: (integer) number of steps to take when optimizing
174     - batch_size: (integer) number of training examples to use at each step.
175     - verbose: (boolean) If true, print progress during optimization.
176
177     Outputs:
178     A list containing the value of the loss function at each training iteration.
179     """
180     num_train, dim = X.shape
181     num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of classes
182
183     self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of self.W
184
185     # Run stochastic gradient descent to optimize W
186     loss_history = []
187
188     for it in np.arange(num_iters):
189         X_batch = None
190         y_batch = None
191
192         # ===== #
193         # YOUR CODE HERE:
194         # Sample batch_size elements from the training data for use in
195         # gradient descent. After sampling,
196         # - X_batch should have shape: (dim, batch_size)
197         # - y_batch should have shape: (batch_size,)
198         # The indices should be randomly generated to reduce correlations
199         # in the dataset. Use np.random.choice. It's okay to sample with
200         # replacement.
201         # ===== #
202         indic = np.random.choice(num_train, batch_size)
203         X_batch = X[indic,:]
204         y_batch = y[indic]
205         # ===== #
206         # END YOUR CODE HERE
207         # ===== #
208
209         # evaluate loss and gradient
210         loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
211         loss_history.append(loss)
212
213         # ===== #
214         # YOUR CODE HERE:
215         # Update the parameters, self.W, with a gradient step
216         # ===== #
217         self.W -= learning_rate*grad
218         # ===== #

```

```
219     # END YOUR CODE HERE
220     # ===== #
221
222     if verbose and it % 100 == 0:
223         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
224
225     return loss_history
226
227 def predict(self, X):
228     """
229     Inputs:
230     - X: N x D array of training data. Each row is a D-dimensional point.
231
232     Returns:
233     - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
234       array of length N, and each element is an integer giving the predicted
235       class.
236     """
237     y_pred = np.zeros(X.shape[0])
238
239
240     # ===== #
241     # YOUR CODE HERE:
242     #   Predict the labels given the training data with the parameter self.W.
243     # ===== #
244     prod = np.dot(X, np.transpose(self.W))
245
246     y_pred = np.argmax(prod, axis = 1)
247     # ===== #
248     # END YOUR CODE HERE
249     # ===== #
250
251     return y_pred
```