

```

1 import numpy as np
2
3 from .layers import *
4 from .layer_utils import *
5
6 """
7 This code was originally written for CS 231n at Stanford University
8 (cs231n.stanford.edu). It has been modified in various areas for use in the
9 ECE 239AS class at UCLA. This includes the descriptions of what code to
10 implement as well as some slight potential changes in variable names to be
11 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
12 permission to use this code. To see the original version, please visit
13 cs231n.stanford.edu.
14 """
15
16 class TwoLayerNet(object):
17     """
18     A two-layer fully-connected neural network with ReLU nonlinearity and
19     softmax loss that uses a modular layer design. We assume an input dimension
20     of D, a hidden dimension of H, and perform classification over C classes.
21
22     The architecture should be affine - relu - affine - softmax.
23
24     Note that this class does not implement gradient descent; instead, it
25     will interact with a separate Solver object that is responsible for running
26     optimization.
27
28     The learnable parameters of the model are stored in the dictionary
29     self.params that maps parameter names to numpy arrays.
30     """
31
32     def __init__(self, input_dim=3*32*32, hidden_dims=100, num_classes=10,
33                 dropout=0, weight_scale=1e-3, reg=0.0):
34         """
35         Initialize a new network.
36
37         Inputs:
38         - input_dim: An integer giving the size of the input
39         - hidden_dims: An integer giving the size of the hidden layer
40         - num_classes: An integer giving the number of classes to classify
41         - dropout: Scalar between 0 and 1 giving dropout strength.
42         - weight_scale: Scalar giving the standard deviation for random
43           initialization of the weights.
44         - reg: Scalar giving L2 regularization strength.
45         """
46         self.params = {}
47         self.reg = reg
48         self.cache = {}
49
50         # ===== #
51         # YOUR CODE HERE:
52         # Initialize W1, W2, b1, and b2. Store these as self.params['W1'],
53         # self.params['W2'], self.params['b1'] and self.params['b2']. The
54         # biases are initialized to zero and the weights are initialized
55         # so that each parameter has mean 0 and standard deviation weight_scale.
56         # The dimensions of W1 should be (input_dim, hidden_dim) and the
57         # dimensions of W2 should be (hidden_dims, num_classes)
58         # ===== #
59
60         np.random.seed(0)
61         self.params = {}
62         self.params['W1'] = weight_scale * np.random.randn(input_dim, hidden_dims)
63         self.params['b1'] = np.zeros(hidden_dims)
64         self.params['W2'] = weight_scale * np.random.randn(hidden_dims, num_classes)
65         self.params['b2'] = np.zeros(num_classes)
66
67         # ===== #
68         # END YOUR CODE HERE
69         # ===== #
70

```

```

71 def loss(self, X, y=None):
72     """
73     Compute loss and gradient for a minibatch of data.
74
75     Inputs:
76     - X: Array of input data of shape (N, d_1, ..., d_k)
77     - y: Array of labels, of shape (N,). y[i] gives the label for X[i].
78
79     Returns:
80     If y is None, then run a test-time forward pass of the model and return:
81     - scores: Array of shape (N, C) giving classification scores, where
82       scores[i, c] is the classification score for X[i] and class c.
83
84     If y is not None, then run a training-time forward and backward pass and
85     return a tuple of:
86     - loss: Scalar value giving the loss
87     - grads: Dictionary with the same keys as self.params, mapping parameter
88       names to gradients of the loss with respect to those parameters.
89     """
90     scores = None
91     W1, b1 = self.params['W1'], self.params['b1']
92     W2, b2 = self.params['W2'], self.params['b2']
93     N = X.shape[0]
94     D = np.prod(X.shape[1:])
95     # ===== #
96     # YOUR CODE HERE:
97     # Implement the forward pass of the two-layer neural network. Store
98     # the class scores as the variable 'scores'. Be sure to use the layers
99     # you prior implemented.
100    # ===== #
101
102    out1, cache1 = affine_relu_forward(X, W1, b1)
103
104    out2, cache2 = affine_forward(out1, W2, b2)
105    scores = out2
106
107    # ===== #
108    # END YOUR CODE HERE
109    # ===== #
110
111    # If y is None then we are in test mode so just return scores
112    if y is None:
113        return scores
114
115    loss, grads = 0, {}
116    # ===== #
117    # YOUR CODE HERE:
118    # Implement the backward pass of the two-layer neural net. Store
119    # the loss as the variable 'loss' and store the gradients in the
120    # 'grads' dictionary. For the grads dictionary, grads['W1'] holds
121    # the gradient for W1, grads['b1'] holds the gradient for b1, etc.
122    # i.e., grads[k] holds the gradient for self.params[k].
123    #
124    # Add L2 regularization, where there is an added cost 0.5*self.reg*W^2
125    # for each W. Be sure to include the 0.5 multiplying factor to
126    # match our implementation.
127    #
128    # And be sure to use the layers you prior implemented.
129    # ===== #
130
131    loss, dscore = softmax_loss(out2, y)
132    loss += 0.5 * self.reg * np.sum(W1*W1) + 0.5*self.reg * np.sum(W2*W2)
133
134    dx1, grads['W2'], grads['b2'] = affine_backward(dscore, cache2)
135    _, grads['W1'], grads['b1'] = affine_relu_backward(dx1, cache1)
136
137
138    grads['W2'] += self.reg * W2
139    grads['W1'] += self.reg * W1
140    # ===== #
141    # END YOUR CODE HERE

```

```

142 # ===== #
143
144 return loss, grads
145
146
147 class FullyConnectedNet(object):
148     """
149     A fully-connected neural network with an arbitrary number of hidden layers,
150     ReLU nonlinearities, and a softmax loss function. This will also implement
151     dropout and batch normalization as options. For a network with L layers,
152     the architecture will be
153
154     {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax
155
156     where batch normalization and dropout are optional, and the {...} block is
157     repeated L - 1 times.
158
159     Similar to the TwoLayerNet above, learnable parameters are stored in the
160     self.params dictionary and will be learned using the Solver class.
161     """
162
163     def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
164                 dropout=0, use_batchnorm=False, reg=0.0,
165                 weight_scale=1e-2, dtype=np.float32, seed=None):
166         """
167         Initialize a new FullyConnectedNet.
168
169         Inputs:
170         - hidden_dims: A list of integers giving the size of each hidden layer.
171         - input_dim: An integer giving the size of the input.
172         - num_classes: An integer giving the number of classes to classify.
173         - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
174           the network should not use dropout at all.
175         - use_batchnorm: Whether or not the network should use batch normalization.
176         - reg: Scalar giving L2 regularization strength.
177         - weight_scale: Scalar giving the standard deviation for random
178           initialization of the weights.
179         - dtype: A numpy datatype object; all computations will be performed using
180           this datatype. float32 is faster but less accurate, so you should use
181           float64 for numeric gradient checking.
182         - seed: If not None, then pass this random seed to the dropout layers. This
183           will make the dropout layers deterministic so we can gradient check the
184           model.
185         """
186         self.use_batchnorm = use_batchnorm
187         self.use_dropout = dropout > 0
188         self.reg = reg
189         self.num_layers = 1 + len(hidden_dims)
190         self.dtype = dtype
191         self.params = {}
192
193         # ===== #
194         # YOUR CODE HERE:
195         # Initialize all parameters of the network in the self.params dictionary.
196         # The weights and biases of layer 1 are W1 and b1; and in general the
197         # weights and biases of layer i are Wi and bi. The
198         # biases are initialized to zero and the weights are initialized
199         # so that each parameter has mean 0 and standard deviation weight_scale.
200         # ===== #
201         for i in range(self.num_layers):
202             if i == 0:
203                 # weights between input and hidden layer
204                 self.params['W' + str(i+1)] = weight_scale * np.random.randn(input_dim, hidden_dims[i])
205                 self.params['b' + str(i+1)] = np.zeros(hidden_dims[i])
206
207                 # weights between hidden and hidden layer
208             elif i < self.num_layers - 1:
209                 self.params['W' + str(i+1)] = weight_scale * np.random.randn(hidden_dims[i-1], hidden_dims[i])
210                 self.params['b' + str(i+1)] = np.zeros(hidden_dims[i])
211
212             # weights between hidden and output layer

```

```

213     else:
214         self.params['W' + str(i+1)] = weight_scale * np.random.randn(hidden_dims[i-1], num_classes)
215         self.params['b' + str(i+1)] = np.zeros(num_classes)
216
217
218     # ===== #
219     # END YOUR CODE HERE
220     # ===== #
221
222     # When using dropout we need to pass a dropout_param dictionary to each
223     # dropout layer so that the layer knows the dropout probability and the mode
224     # (train / test). You can pass the same dropout_param to each dropout layer.
225     self.dropout_param = {}
226     if self.use_dropout:
227         self.dropout_param = {'mode': 'train', 'p': dropout}
228         if seed is not None:
229             self.dropout_param['seed'] = seed
230
231     # With batch normalization we need to keep track of running means and
232     # variances, so we need to pass a special bn_param object to each batch
233     # normalization layer. You should pass self.bn_params[0] to the forward pass
234     # of the first batch normalization layer, self.bn_params[1] to the forward
235     # pass of the second batch normalization layer, etc.
236     self.bn_params = []
237     if self.use_batchnorm:
238         self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]
239
240     # Cast all parameters to the correct datatype
241     for k, v in self.params.items():
242         self.params[k] = v.astype(dtype)
243
244
245 def loss(self, X, y=None):
246     """
247     Compute loss and gradient for the fully-connected net.
248
249     Input / output: Same as TwoLayerNet above.
250     """
251     X = X.astype(self.dtype)
252     mode = 'test' if y is None else 'train'
253
254     # Set train/test mode for batchnorm params and dropout param since they
255     # behave differently during training and testing.
256     if self.dropout_param is not None:
257         self.dropout_param['mode'] = mode
258     if self.use_batchnorm:
259         for bn_param in self.bn_params:
260             bn_param['mode'] = mode
261
262     scores = None
263     caches = []
264     # ===== #
265     # YOUR CODE HERE:
266     # Implement the forward pass of the FC net and store the output
267     # scores as the variable "scores".
268     # ===== #
269
270     data_in = X
271     for i in range(self.num_layers):
272
273         if i < self.num_layers - 1:
274             out, cache = affine_relu_forward(data_in, self.params['W'+str(i+1)], self.params['b'+str(i+1)])
275             data_in = out
276             caches.append(cache)
277         else:
278             out2, cache2 = affine_forward(data_in, self.params['W'+str(i+1)], self.params['b'+str(i+1)])
279
280     scores = out2
281
282
283     # ===== #

```

```

284 # END YOUR CODE HERE
285 # ===== #
286
287 # If test mode return early
288 if mode == 'test':
289     return scores
290
291 loss, grads = 0.0, {}
292 # ===== #
293 # YOUR CODE HERE:
294 # Implement the backwards pass of the FC net and store the gradients
295 # in the grads dict, so that grads[k] is the gradient of self.params[k]
296 # Be sure your L2 regularization includes a 0.5 factor.
297 # ===== #
298
299 loss, dscore = softmax_loss(out2, y)
300 regularization = 0
301 for i in range(self.num_layers):
302     #regularization += 0.5 * self.reg * np.sum(self.params['W'+str(i+1)])
303     loss += 0.5 * self.reg * np.sum(self.params['W'+str(i+1)] * self.params['W'+str(i+1)])
304
305 for i in reversed(range(self.num_layers)):
306     if i == self.num_layers - 1:
307         dx1, grads['W'+str(i+1)], grads['b'+str(i+1)] = affine_backward(dscore, cache2)
308     else:
309         dx, grads['W'+str(i+1)], grads['b'+str(i+1)] = affine_relu_backward(dx1, caches[i])
310         dx1 = dx
311
312 for i in reversed(range(self.num_layers)):
313     grads['W'+str(i+1)] += self.reg * self.params['W'+str(i+1)]
314
315
316 # ===== #
317 # END YOUR CODE HERE
318 # ===== #
319 return loss, grads

```