

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 """
5 This code was originally written for CS 231n at Stanford University
6 (cs231n.stanford.edu). It has been modified in various areas for use in the
7 ECE 239AS class at UCLA. This includes the descriptions of what code to
8 implement as well as some slight potential changes in variable names to be
9 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
10 permission to use this code. To see the original version, please visit
11 cs231n.stanford.edu.
12 """
13
14 class TwoLayerNet(object):
15     """
16     A two-layer fully-connected neural network. The net has an input dimension of
17     N, a hidden layer dimension of H, and performs classification over C classes.
18     We train the network with a softmax loss function and L2 regularization on the
19     weight matrices. The network uses a ReLU nonlinearity after the first fully
20     connected layer.
21
22     In other words, the network has the following architecture:
23
24     input - fully connected layer - ReLU - fully connected layer - softmax
25
26     The outputs of the second fully-connected layer are the scores for each class.
27     """
28
29     def __init__(self, input_size, hidden_size, output_size, std=1e-4):
30         """
31         Initialize the model. Weights are initialized to small random values and
32         biases are initialized to zero. Weights and biases are stored in the
33         variable self.params, which is a dictionary with the following keys:
34
35         W1: First layer weights; has shape (H, D)
36         b1: First layer biases; has shape (H,)
37         W2: Second layer weights; has shape (C, H)
38         b2: Second layer biases; has shape (C,)
39
40         Inputs:
41         - input_size: The dimension D of the input data.
42         - hidden_size: The number of neurons H in the hidden layer.
43         - output_size: The number of classes C.
44         """
45         np.random.seed(0)
46         self.params = {}
47         self.params['W1'] = std * np.random.randn(hidden_size, input_size)
48         self.params['b1'] = np.zeros(hidden_size)
49         self.params['W2'] = std * np.random.randn(output_size, hidden_size)
50         self.params['b2'] = np.zeros(output_size)
51
52     def loss(self, X, y=None, reg=0.0):
53         """
54         Compute the loss and gradients for a two layer fully connected neural
55         network.
56
57         Inputs:
58         - X: Input data of shape (N, D). Each X[i] is a training sample.
59         - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
60             an integer in the range 0 <= y[i] < C. This parameter is optional; if it
61             is not passed then we only return scores, and if it is passed then we
62             instead return the loss and gradients.
63         - reg: Regularization strength.
64
65         Returns:
66         If y is None, return a matrix scores of shape (N, C) where scores[i, c] is
67         the score for class c on input X[i].
68         """

```

```

70 If y is not None, instead return a tuple of:
71 - loss: Loss (data loss and regularization loss) for this batch of training
72 samples.
73 - grads: Dictionary mapping parameter names to gradients of those parameters
74 with respect to the loss function; has the same keys as self.params.
75 """
76 # Unpack variables from the params dictionary
77 W1, b1 = self.params['W1'], self.params['b1']
78 W2, b2 = self.params['W2'], self.params['b2']
79 N, D = X.shape
80
81 # Compute the forward pass
82 scores = None
83 # ===== #
84 # YOUR CODE HERE:
85 # Calculate the output scores of the neural network. The result
86 # should be (C, N). As stated in the description for this class,
87 # there should not be a ReLU layer after the second FC layer.
88 # The output of the second FC layer is the output scores. Do not
89 # use a for loop in your implementation.
90 # ===== #
91
92 h1 = np.maximum(0, np.dot(W1, X.T) + np.matrix(b1).T)
93 z = np.dot(W2, h1) + np.matrix(b2).T
94 scores = z.T
95 pass
96
97 # ===== #
98 # END YOUR CODE HERE
99 # ===== #
100 softmax = lambda x: np.exp(x - np.max(x)) / np.exp(x - np.max(x)).sum(axis=0)
101 # If the targets are not given then jump out, we're done
102 if y is None:
103     return scores
104
105 # Compute the loss
106 loss = None
107
108 # ===== #
109 # YOUR CODE HERE:
110 # Calculate the loss of the neural network. This includes the
111 # softmax loss and the L2 regularization for W1 and W2. Store the
112 # total loss in the variable loss. Multiply the regularization
113 # loss by 0.5 (in addition to the factor reg).
114 # ===== #
115
116 # scores is num_examples by num_classes
117 aExp = np.exp(scores)
118
119 prob = aExp / np.sum(aExp, axis = 1)
120 correctLogProb = -np.log(prob[range(N), y])
121 dataLoss = np.sum(correctLogProb) / N
122 regloss = 0.5 * reg * np.sum(W1 * W1) + 0.5 * reg * np.sum(W2 * W2)
123
124 loss = regloss + dataLoss
125 # ===== #
126 # END YOUR CODE HERE
127 # ===== #
128
129 grads = {}
130
131 # ===== #
132 # YOUR CODE HERE:
133 # Implement the backward pass. Compute the derivatives of the
134 # weights and the biases. Store the results in the grads
135 # dictionary. e.g., grads['W1'] should store the gradient for
136 # W1, and be of the same size as W1.
137 # ===== #
138 # prob = np.matrix(prob).T
139 # print y.shape[0]

```

```

140 prob = np.arange(y.shape[0], y) - 1
141 prob /= y.shape[0]
142 grads['W2'] = np.dot(h1, prob).T
143 grads['b2'] = np.sum(prob, axis=0)
144
145 dh1 = np.dot(prob, W2).T
146 dh1[h1<=0] = 0
147
148 grads['W1'] = np.dot(X.T, dh1.T).T
149 grads['b1'] = np.sum(dh1, axis=1).T
150
151
152
153 grads['W1'] += reg*W1
154 grads['W2'] += reg*W2
155
156
157 # ===== #
158 # END YOUR CODE HERE
159 # ===== #
160
161 return loss, grads
162
163 def train(self, X, y, X_val, y_val,
164         learning_rate=1e-3, learning_rate_decay=0.95,
165         reg=1e-5, num_iters=100,
166         batch_size=200, verbose=False):
167     """
168     Train this neural network using stochastic gradient descent.
169
170     Inputs:
171     - X: A numpy array of shape (N, D) giving training data.
172     - y: A numpy array of shape (N,) giving training labels; y[i] = c means that
173         X[i] has label c, where 0 <= c < C.
174     - X_val: A numpy array of shape (N_val, D) giving validation data.
175     - y_val: A numpy array of shape (N_val,) giving validation labels.
176     - learning_rate: Scalar giving learning rate for optimization.
177     - learning_rate_decay: Scalar giving factor used to decay the learning rate
178         after each epoch.
179     - reg: Scalar giving regularization strength.
180     - num_iters: Number of steps to take when optimizing.
181     - batch_size: Number of training examples to use per step.
182     - verbose: boolean; if true print progress during optimization.
183     """
184     num_train = X.shape[0]
185     iterations_per_epoch = max(num_train / batch_size, 1)
186
187     # Use SGD to optimize the parameters in self.model
188     loss_history = []
189     train_acc_history = []
190     val_acc_history = []
191
192     for it in np.arange(num_iters):
193         X_batch = None
194         y_batch = None
195
196         # ===== #
197         # YOUR CODE HERE:
198         # Create a minibatch by sampling batch_size samples randomly.
199         # ===== #
200         indic = np.random.choice(num_train, batch_size)
201         X_batch = X[indic,:]
202         y_batch = y[indic]
203
204         # ===== #
205         # END YOUR CODE HERE
206         # ===== #
207
208         # Compute loss and gradients using the current minibatch
209         loss, grads = self.loss(X_batch, y=y_batch, reg=reg)

```

```

210     loss_history.append(loss)
211
212     # ===== #
213     # YOUR CODE HERE:
214     #     Perform a gradient descent step using the minibatch to update
215     #     all parameters (i.e., W1, W2, b1, and b2).
216     # ===== #
217     self.params['W1'] -= learning_rate*grads['W1']
218     self.params['W2'] -= learning_rate*grads['W2']
219     self.params['b1'] -= learning_rate*np.asarray(grads['b1']).reshape(-1)
220     self.params['b2'] -= learning_rate*np.asarray(grads['b2']).reshape(-1)
221
222
223     # ===== #
224     # END YOUR CODE HERE
225     # ===== #
226
227     if verbose and it % 100 == 0:
228         print('iteration {} / {}: loss {}'.format(it, num_iters, loss))
229
230     # Every epoch, check train and val accuracy and decay learning rate.
231     if it % iterations_per_epoch == 0:
232         # Check accuracy
233         train_acc = (self.predict(X_batch) == y_batch).mean()
234         val_acc = (self.predict(X_val) == y_val).mean()
235         train_acc_history.append(train_acc)
236         val_acc_history.append(val_acc)
237
238         # Decay learning rate
239         learning_rate *= learning_rate_decay
240
241     return {
242         'loss_history': loss_history,
243         'train_acc_history': train_acc_history,
244         'val_acc_history': val_acc_history,
245     }
246
247 def predict(self, X):
248     """
249     Use the trained weights of this two-layer network to predict labels for
250     data points. For each data point we predict scores for each of the C
251     classes, and assign each data point to the class with the highest score.
252
253     Inputs:
254     - X: A numpy array of shape (N, D) giving N D-dimensional data points to
255         classify.
256
257     Returns:
258     - y_pred: A numpy array of shape (N,) giving predicted labels for each of
259         the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
260         to have class c, where 0 <= c < C.
261     """
262     y_pred = None
263
264     # ===== #
265     # YOUR CODE HERE:
266     #     Predict the class given the input data.
267     # ===== #
268     #reluFunc = lambda x: np.multiply(x,(x>0))
269     z = np.dot(X, self.params['W1'].T) + self.params['b1']
270     h1 = np.maximum(0, z)
271     out = np.dot(h1, self.params['W2'].T) + self.params['b2']
272
273     y_pred = np.argmax(out, axis = 1)
274
275
276     # ===== #
277     # END YOUR CODE HERE
278     # ===== #

```

279

280 **return** y_pred