

Convolutional neural network layers

In this notebook, we will build the convolutional neural network layers. This will be followed by a spatial batchnorm, and then in the final notebook of this assignment, we will train a CNN to further improve the validation accuracy on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

```
In [2]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The autoreload extension is already loaded. To reload it, use:
`%reload_ext autoreload`

Implementing CNN layers

Just as we implemented modular layers for fully connected networks, batch normalization, and dropout, we'll want to implement modular layers for convolutional neural networks. These layers are in `nndl/conv_layers.py`.

Convolutional forward pass

Begin by implementing a naive version of the forward pass of the CNN that uses `for` loops. This function is `conv_forward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a triple `for` loop.

After you implement `conv_forward_naive`, test your implementation by running the cell below.

```
In [5]: x_shape = (2, 3, 4, 4)
w_shape = (3, 3, 4, 4)
x = np.linspace(-0.1, 0.5, num=np.prod(x_shape)).reshape(x_shape)
w = np.linspace(-0.2, 0.3, num=np.prod(w_shape)).reshape(w_shape)
b = np.linspace(-0.1, 0.2, num=3)

conv_param = {'stride': 2, 'pad': 1}
out, _ = conv_forward_naive(x, w, b, conv_param)
correct_out = np.array([[[[-0.08759809, -0.10987781],
                           [-0.18387192, -0.2109216 ]],
                          [[ 0.21027089,  0.21661097],
                           [ 0.22847626,  0.23004637]],
                          [[ 0.50813986,  0.54309974],
                           [ 0.64082444,  0.67101435]]]],
                        [[[-0.98053589, -1.03143541],
                           [-1.19128892, -1.24695841]],
                          [[ 0.69108355,  0.66880383],
                           [ 0.59480972,  0.56776003]],
                          [[ 2.36270298,  2.36904306],
                           [ 2.38090835,  2.38247847]]]])

# Compare your output to ours; difference should be around 1e-8
print('Testing conv_forward_naive')
print('difference: ', rel_error(out, correct_out))

Testing conv_forward_naive
('difference: ', 2.2121476417505994e-08)
```

Convolutional backward pass

Now, implement a naive version of the backward pass of the CNN. The function is `conv_backward_naive` in `nndl/conv_layers.py`. Don't worry about efficiency of implementation. Later on, we provide a fast implementation of these layers. This version ought to test your understanding of convolution. In our implementation, there is a quadruple `for` loop.

After you implement `conv_backward_naive`, test your implementation by running the cell below.

```

In [11]: x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_forward_naive(x,w,b,conv_param)

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_naive(x, w
, b, conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_naive(x, w
, b, conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_forward_naive(x, w
, b, conv_param)[0], b, dout)

out, cache = conv_forward_naive(x, w, b, conv_param)
dx, dw, db = conv_backward_naive(dout, cache)

# Your errors should be around 1e-9
print('Testing conv_backward_naive function')
print('dx error: ', rel_error(dx, dx_num))
print('dw error: ', rel_error(dw, dw_num))
print('db error: ', rel_error(db, db_num))

Testing conv_backward_naive function
('dx error: ', 9.52394925623475e-09)
('dw error: ', 9.704958923846182e-10)
('db error: ', 3.7489601082852155e-11)

```

Max pool forward pass

In this section, we will implement the forward pass of the max pool. The function is `max_pool_forward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_forward_naive`, test your implementation by running the cell below.

```
In [19]: x_shape = (2, 3, 4, 4)
x = np.linspace(-0.3, 0.4, num=np.prod(x_shape)).reshape(x_shape)
pool_param = {'pool_width': 2, 'pool_height': 2, 'stride': 2}

out, _ = max_pool_forward_naive(x, pool_param)

correct_out = np.array([[[[-0.26315789, -0.24842105],
                           [-0.20421053, -0.18947368]],
                          [[-0.14526316, -0.13052632],
                           [-0.08631579, -0.07157895]],
                          [[-0.02736842, -0.01263158],
                           [ 0.03157895,  0.04631579]]],
                        [[[ 0.09052632,  0.10526316],
                           [ 0.14947368,  0.16421053]],
                          [[ 0.20842105,  0.22315789],
                           [ 0.26736842,  0.28210526]],
                          [[ 0.32631579,  0.34105263],
                           [ 0.38526316,  0.4          ]]]]])

# Compare your output with ours. Difference should be around 1e-8.
print('Testing max_pool_forward_naive function:')
print('difference: ', rel_error(out, correct_out))
```

```
Testing max_pool_forward_naive function:
('difference: ', 4.1666665157267834e-08)
```

Max pool backward pass

In this section, you will implement the backward pass of the max pool. The function is `max_pool_backward_naive` in `nndl/conv_layers.py`. Do not worry about the efficiency of implementation.

After you implement `max_pool_backward_naive`, test your implementation by running the cell below.

```
In [21]: x = np.random.randn(3, 2, 8, 8)
dout = np.random.randn(3, 2, 4, 4)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

dx_num = eval_numerical_gradient_array(lambda x: max_pool_forward_naive(
x, pool_param)[0], x, dout)

out, cache = max_pool_forward_naive(x, pool_param)
dx = max_pool_backward_naive(dout, cache)

# Your error should be around 1e-12
print('Testing max_pool_backward_naive function:')
print('dx error: ', rel_error(dx, dx_num))
```

```
Testing max_pool_backward_naive function:
('dx error: ', 3.2756263854194794e-12)
```

Fast implementation of the CNN layers

Implementing fast versions of the CNN layers can be difficult. We will provide you with the fast layers implemented by cs231n. They are provided in `cs231n/fast_layers.py`.

The fast convolution implementation depends on a Cython extension; to compile it you need to run the following from the `cs231n` directory:

```
python setup.py build_ext --inplace
```

NOTE: The fast implementation for pooling will only perform optimally if the pooling regions are non-overlapping and tile the input. If these conditions are not met then the fast pooling implementation will not be much faster than the naive implementation.

You can compare the performance of the naive and fast versions of these layers by running the cell below.

You should see pretty drastic speedups in the implementation of these layers. On our machine, the forward pass speeds up by 17x and the backward pass speeds up by 840x. Of course, these numbers will vary from machine to machine, as well as on your precise implementation of the naive layers.

```

In [22]: from cs231n.fast_layers import conv_forward_fast, conv_backward_fast
         from time import time

x = np.random.randn(100, 3, 31, 31)
w = np.random.randn(25, 3, 3, 3)
b = np.random.randn(25,)
dout = np.random.randn(100, 25, 16, 16)
conv_param = {'stride': 2, 'pad': 1}

t0 = time()
out_naive, cache_naive = conv_forward_naive(x, w, b, conv_param)
t1 = time()
out_fast, cache_fast = conv_forward_fast(x, w, b, conv_param)
t2 = time()

print('Testing conv_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('Difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive, dw_naive, db_naive = conv_backward_naive(dout, cache_naive)
t1 = time()
dx_fast, dw_fast, db_fast = conv_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting conv_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('Fast: %fs' % (t2 - t1))
print('Speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))
print('dw difference: ', rel_error(dw_naive, dw_fast))
print('db difference: ', rel_error(db_naive, db_fast))

Testing conv_forward_fast:
Naive: 5.832725s
Fast: 0.022420s
Speedup: 260.155279x
('Difference: ', 2.879853592390352e-10)

Testing conv_backward_fast:
Naive: 9.468566s
Fast: 0.012800s
Speedup: 739.732971x
('dx difference: ', 1.886313090589319e-11)
('dw difference: ', 1.399482885637161e-12)
('db difference: ', 1.015514672183838e-13)

```

```
In [23]: from cs231n.fast_layers import max_pool_forward_fast, max_pool_backward_fast

x = np.random.randn(100, 3, 32, 32)
dout = np.random.randn(100, 3, 16, 16)
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

t0 = time()
out_naive, cache_naive = max_pool_forward_naive(x, pool_param)
t1 = time()
out_fast, cache_fast = max_pool_forward_fast(x, pool_param)
t2 = time()

print('Testing pool_forward_fast:')
print('Naive: %fs' % (t1 - t0))
print('fast: %fs' % (t2 - t1))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('difference: ', rel_error(out_naive, out_fast))

t0 = time()
dx_naive = max_pool_backward_naive(dout, cache_naive)
t1 = time()
dx_fast = max_pool_backward_fast(dout, cache_fast)
t2 = time()

print('\nTesting pool_backward_fast:')
print('Naive: %fs' % (t1 - t0))
print('speedup: %fx' % ((t1 - t0) / (t2 - t1)))
print('dx difference: ', rel_error(dx_naive, dx_fast))

Testing pool_forward_fast:
Naive: 0.415331s
fast: 0.004256s
speedup: 97.586970x
('difference: ', 0.0)

Testing pool_backward_fast:
Naive: 0.995473s
speedup: 56.622145x
('dx difference: ', 0.0)
```

Implementation of cascaded layers

We've provided the following functions in `nndl/conv_layer_utils.py`:

- `conv_relu_forward`
- `conv_relu_backward`
- `conv_relu_pool_forward`
- `conv_relu_pool_backward`

These use the fast implementations of the conv net layers. You can test them below:

```
In [24]: from nndl.conv_layer_utils import conv_relu_pool_forward, conv_relu_pool_
        _backward

x = np.random.randn(2, 3, 16, 16)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}
pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}

out, cache = conv_relu_pool_forward(x, w, b, conv_param, pool_param)
dx, dw, db = conv_relu_pool_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_pool_forward(
x, w, b, conv_param, pool_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_pool_forward(
x, w, b, conv_param, pool_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_pool_forward(
x, w, b, conv_param, pool_param)[0], b, dout)

print('Testing conv_relu_pool')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))
```

```
Testing conv_relu_pool
('dx error: ', 7.680745635579335e-09)
('dw error: ', 1.2848793336321723e-10)
('db error: ', 7.103094323592727e-11)
```



```
In [25]: from nndl.conv_layer_utils import conv_relu_forward, conv_relu_backward

x = np.random.randn(2, 3, 8, 8)
w = np.random.randn(3, 3, 3, 3)
b = np.random.randn(3,)
dout = np.random.randn(2, 3, 8, 8)
conv_param = {'stride': 1, 'pad': 1}

out, cache = conv_relu_forward(x, w, b, conv_param)
dx, dw, db = conv_relu_backward(dout, cache)

dx_num = eval_numerical_gradient_array(lambda x: conv_relu_forward(x, w,
    b, conv_param)[0], x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_relu_forward(x, w,
    b, conv_param)[0], w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_relu_forward(x, w,
    b, conv_param)[0], b, dout)

print('Testing conv_relu:')
print('dx error: ', rel_error(dx_num, dx))
print('dw error: ', rel_error(dw_num, dw))
print('db error: ', rel_error(db_num, db))

Testing conv_relu:
('dx error: ', 2.455290779970179e-09)
('dw error: ', 2.312204258373735e-09)
('db error: ', 1.5098341437876588e-11)
```

What next?

We saw how helpful batch normalization was for training FC nets. In the next notebook, we'll implement a batch normalization for convolutional neural networks, and then finish off by implementing a CNN to improve our validation accuracy on CIFAR-10.

Spatial batch normalization

In fully connected networks, we performed batch normalization on the activations. To do something equivalent on CNNs, we modify batch normalization slightly.

Normally batch-normalization accepts inputs of shape (N, D) and produces outputs of shape (N, D) , where we normalize across the minibatch dimension N . For data coming from convolutional layers, batch normalization accepts inputs of shape (N, C, H, W) and produces outputs of shape (N, C, H, W) where the N dimension gives the minibatch size and the (H, W) dimensions give the spatial size of the feature map.

How do we calculate the spatial averages? First, notice that for the C feature maps we have (i.e., the layer has C filters) that each of these ought to have its own batch norm statistics, since each feature map may be picking out very different features in the images. However, within a feature map, we may assume that across all inputs and across all locations in the feature map, there ought to be relatively similar first and second order statistics. Hence, one way to think of spatial batch-normalization is to reshape the (N, C, H, W) array as an $(N \cdot H \cdot W, C)$ array and perform batch normalization on this array.

Since spatial batch norm and batch normalization are similar, it'd be good to at this point also copy and paste our prior implemented layers from HW #4. Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4:

- `layers.py` for your FC network layers, as well as `batchnorm` and `dropout`.
- `layer_utils.py` for your combined FC network layers.
- `optim.py` for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

If you use your prior implementations of the `batchnorm`, then your spatial `batchnorm` implementation may be very short. Our implementations of the forward and backward pass are each 6 lines of code.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

```
In [1]: ## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.conv_layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_array
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

Spatial batch normalization forward pass

Implement the forward pass, `spatial_batchnorm_forward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```
In [15]: # Check the training-time forward pass by checking means and variances
# of features both before and after spatial batch normalization

N, C, H, W = 2, 3, 4, 5
x = 4 * np.random.randn(N, C, H, W) + 10

print('Before spatial batch normalization:')
print('  Shape: ', x.shape)
print('  Means: ', x.mean(axis=(0, 2, 3)))
print('  Stds: ', x.std(axis=(0, 2, 3)))

# Means should be close to zero and stds close to one
gamma, beta = np.ones(C), np.zeros(C)
bn_param = {'mode': 'train'}
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization:')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))

# Means should be close to beta and stds close to gamma
gamma, beta = np.asarray([3, 4, 5]), np.asarray([6, 7, 8])
out, _ = spatial_batchnorm_forward(x, gamma, beta, bn_param)
print('After spatial batch normalization (nontrivial gamma, beta):')
print('  Shape: ', out.shape)
print('  Means: ', out.mean(axis=(0, 2, 3)))
print('  Stds: ', out.std(axis=(0, 2, 3)))
```

```
Before spatial batch normalization:
('  Shape: ', (2, 3, 4, 5))
('  Means: ', array([10.11007759,  9.35884986, 10.14078303]))
('  Stds: ', array([3.64895311, 3.80273697, 2.25507063]))
After spatial batch normalization:
('  Shape: ', (2, 3, 4, 5))
('  Means: ', array([ 2.14064877e-16, -4.57966998e-16, -1.80411242e-16]))
('  Stds: ', array([0.99999962, 0.99999965, 0.99999902]))
After spatial batch normalization (nontrivial gamma, beta):
('  Shape: ', (2, 3, 4, 5))
('  Means: ', array([6., 7., 8.]))
('  Stds: ', array([2.99999887, 3.99999862, 4.99999508]))
```

Spatial batch normalization backward pass

Implement the backward pass, `spatial_batchnorm_backward` in `nndl/conv_layers.py`. Test your implementation by running the cell below.

```
In [17]: N, C, H, W = 2, 3, 4, 5
x = 5 * np.random.randn(N, C, H, W) + 12
gamma = np.random.randn(C)
beta = np.random.randn(C)
dout = np.random.randn(N, C, H, W)

bn_param = {'mode': 'train'}
fx = lambda x: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fg = lambda a: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]
fb = lambda b: spatial_batchnorm_forward(x, gamma, beta, bn_param)[0]

dx_num = eval_numerical_gradient_array(fx, x, dout)
da_num = eval_numerical_gradient_array(fg, gamma, dout)
db_num = eval_numerical_gradient_array(fb, beta, dout)

_, cache = spatial_batchnorm_forward(x, gamma, beta, bn_param)
dx, dgamma, dbeta = spatial_batchnorm_backward(dout, cache)
print('dx error: ', rel_error(dx_num, dx))
print('dgamma error: ', rel_error(da_num, dgamma))
print('dbeta error: ', rel_error(db_num, dbeta))

('dx error: ', 6.015817725060696e-08)
('dgamma error: ', 1.3769169556026553e-12)
('dbeta error: ', 5.827940219536585e-12)
```

Convolutional neural networks

In this notebook, we'll put together our convolutional layers to implement a 3-layer CNN. Then, we'll ask you to implement a CNN that can achieve $> 65\%$ validation error on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, their layer structure, and their implementation of fast CNN layers. This also includes `nndl.fc_net`, `nndl.layers`, and `nndl.layer_utils`. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

If you have not completed the Spatial BatchNorm Notebook, please see the following description from that notebook:

Please copy and paste your prior implemented code from HW #4 to start this assignment. If you did not correctly implement the layers in HW #4, you may collaborate with a classmate to use their layer implementations from HW #4. You may also visit TA or Prof OH to correct your implementation.

You'll want to copy and paste from HW #4:

- `layers.py` for your FC network layers, as well as `batchnorm` and `dropout`.
- `layer_utils.py` for your combined FC network layers.
- `optim.py` for your optimizers.

Be sure to place these in the `nndl/` directory so they're imported correctly. Note, as announced in class, we will not be releasing our solutions.

```

In [11]: # As usual, a bit of setup

import numpy as np
import matplotlib.pyplot as plt
from nndl.cnn import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient_array, eval_numerical_gradient
from nndl.layers import *
from nndl.conv_layers import *
from cs231n.fast_layers import *
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

```

The autoreload extension is already loaded. To reload it, use:

```
%reload_ext autoreload
```

```

In [12]: # Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
    print('{}: {}'.format(k, data[k].shape))

X_val: (1000, 3, 32, 32)
X_train: (49000, 3, 32, 32)
X_test: (1000, 3, 32, 32)
y_val: (1000,)
y_train: (49000,)
y_test: (1000,)

```

Three layer CNN

In this notebook, you will implement a three layer CNN. The `ThreeLayerConvNet` class is in `nndl/cnn.py`. You'll need to modify that code for this section, including the initialization, as well as the calculation of the loss and gradients. You should be able to use the building blocks you have either earlier coded or that we have provided. Be sure to use the fast layers.

The architecture of this CNN will be:

conv - relu - 2x2 max pool - affine - relu - affine - softmax

We won't use batchnorm yet. You've also done enough of these to know how to debug; use the cells below.

Note: As we are implementing several layers CNN networks. The gradient error can be expected for the `eval_numerical_gradient()` function. If your W1 max relative error and W2 max relative error are around or below 0.01, they should be acceptable. Other errors should be less than $1e-5$.

```
In [19]: num_inputs = 2
input_dim = (3, 16, 16)
reg = 0.0
num_classes = 10
X = np.random.randn(num_inputs, *input_dim)
y = np.random.randint(num_classes, size=num_inputs)

model = ThreeLayerConvNet(num_filters=3, filter_size=3,
                           input_dim=input_dim, hidden_dim=7,
                           dtype=np.float64)

loss, grads = model.loss(X, y)
for param_name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    param_grad_num = eval_numerical_gradient(f, model.params[param_name],
    verbose=False, h=1e-6)
    e = rel_error(param_grad_num, grads[param_name])
    print('{} max relative error: {}'.format(param_name, rel_error(param_grad_num, grads[param_name])))

W1 max relative error: 0.00048500405658
W2 max relative error: 0.0121570642185
W3 max relative error: 4.34491777159e-05
b1 max relative error: 7.99716890938e-05
b2 max relative error: 2.5868500689e-07
b3 max relative error: 1.14809756429e-09
```

Overfit small dataset

To check your CNN implementation, let's overfit a small dataset.


```
In [20]: num_train = 100
small_data = {
    'X_train': data['X_train'][:num_train],
    'y_train': data['y_train'][:num_train],
    'X_val': data['X_val'],
    'y_val': data['y_val'],
}


model = ThreeLayerConvNet(weight_scale=1e-2)

solver = Solver(model, small_data,
                num_epochs=10, batch_size=50,
                update_rule='adam',
                optim_config={
                    'learning_rate': 1e-3,
                },
                verbose=True, print_every=1)
solver.train()

(Iteration 1 / 20) loss: 2.404106
(Epoch 0 / 10) train acc: 0.230000; val_acc: 0.132000
(Iteration 2 / 20) loss: 2.818953
(Epoch 1 / 10) train acc: 0.220000; val_acc: 0.142000
(Iteration 3 / 20) loss: 2.435576
(Iteration 4 / 20) loss: 2.501954
(Epoch 2 / 10) train acc: 0.240000; val_acc: 0.136000
(Iteration 5 / 20) loss: 2.364800
(Iteration 6 / 20) loss: 1.936613
(Epoch 3 / 10) train acc: 0.230000; val_acc: 0.086000
(Iteration 7 / 20) loss: 1.866935
(Iteration 8 / 20) loss: 1.644003
(Epoch 4 / 10) train acc: 0.530000; val_acc: 0.188000
(Iteration 9 / 20) loss: 1.515180
(Iteration 10 / 20) loss: 1.819292
(Epoch 5 / 10) train acc: 0.580000; val_acc: 0.185000
(Iteration 11 / 20) loss: 1.203726
(Iteration 12 / 20) loss: 1.045983
(Epoch 6 / 10) train acc: 0.600000; val_acc: 0.188000
(Iteration 13 / 20) loss: 1.076013
(Iteration 14 / 20) loss: 0.929915
(Epoch 7 / 10) train acc: 0.760000; val_acc: 0.222000
(Iteration 15 / 20) loss: 0.869404
(Iteration 16 / 20) loss: 0.616786
(Epoch 8 / 10) train acc: 0.750000; val_acc: 0.242000
(Iteration 17 / 20) loss: 0.689910
(Iteration 18 / 20) loss: 0.579885
(Epoch 9 / 10) train acc: 0.820000; val_acc: 0.230000
(Iteration 19 / 20) loss: 0.541249
(Iteration 20 / 20) loss: 0.358143
(Epoch 10 / 10) train acc: 0.840000; val_acc: 0.199000
```

```
In [21]: plt.subplot(2, 1, 1)
plt.plot(solver.loss_history, 'o')
plt.xlabel('iteration')
plt.ylabel('loss')

plt.subplot(2, 1, 2)
plt.plot(solver.train_acc_history, '-o')
plt.plot(solver.val_acc_history, '-o')
plt.legend(['train', 'val'], loc='upper left')
plt.xlabel('epoch')
plt.ylabel('accuracy')
plt.show()
```



Train the network

Now we train the 3 layer CNN on CIFAR-10 and assess its accuracy.

```
In [22]: model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001)

solver = Solver(model, data,
                 num_epochs=1, batch_size=50,
                 update_rule='adam',
                 optim_config={
                     'learning_rate': 1e-3,
                 },
                 verbose=True, print_every=20)
solver.train()
```

```
(Iteration 1 / 980) loss: 2.304829
(Epoch 0 / 1) train acc: 0.110000; val_acc: 0.118000
(Iteration 21 / 980) loss: 2.006605
(Iteration 41 / 980) loss: 2.006921
(Iteration 61 / 980) loss: 2.084292
(Iteration 81 / 980) loss: 1.891646
(Iteration 101 / 980) loss: 1.870840
(Iteration 121 / 980) loss: 1.800872
(Iteration 141 / 980) loss: 1.968251
(Iteration 161 / 980) loss: 1.756411
(Iteration 181 / 980) loss: 1.633845
(Iteration 201 / 980) loss: 1.800655
(Iteration 221 / 980) loss: 1.777673
(Iteration 241 / 980) loss: 1.688106
(Iteration 261 / 980) loss: 1.372005
(Iteration 281 / 980) loss: 1.946737
(Iteration 301 / 980) loss: 1.726557
(Iteration 321 / 980) loss: 1.475561
(Iteration 341 / 980) loss: 1.834388
(Iteration 361 / 980) loss: 1.466202
(Iteration 381 / 980) loss: 1.680460
(Iteration 401 / 980) loss: 1.496826
(Iteration 421 / 980) loss: 1.852106
(Iteration 441 / 980) loss: 1.837103
(Iteration 461 / 980) loss: 1.966308
(Iteration 481 / 980) loss: 1.526592
(Iteration 501 / 980) loss: 1.470095
(Iteration 521 / 980) loss: 1.710456
(Iteration 541 / 980) loss: 1.910397
(Iteration 561 / 980) loss: 1.449640
(Iteration 581 / 980) loss: 1.811990
(Iteration 601 / 980) loss: 1.514811
(Iteration 621 / 980) loss: 1.606383
(Iteration 641 / 980) loss: 1.444682
(Iteration 661 / 980) loss: 1.636750
(Iteration 681 / 980) loss: 1.532761
(Iteration 701 / 980) loss: 1.874972
(Iteration 721 / 980) loss: 1.840662
(Iteration 741 / 980) loss: 1.537102
(Iteration 761 / 980) loss: 1.722641
(Iteration 781 / 980) loss: 1.087972
(Iteration 801 / 980) loss: 1.379641
(Iteration 821 / 980) loss: 1.450924
(Iteration 841 / 980) loss: 1.437804
(Iteration 861 / 980) loss: 1.587830
(Iteration 881 / 980) loss: 1.610635
(Iteration 901 / 980) loss: 1.561900
(Iteration 921 / 980) loss: 1.410366
(Iteration 941 / 980) loss: 1.693163
(Iteration 961 / 980) loss: 1.661325
(Epoch 1 / 1) train acc: 0.461000; val_acc: 0.467000
```

Get > 65% validation accuracy on CIFAR-10.

In the last part of the assignment, we'll now ask you to train a CNN to get better than 65% validation accuracy on CIFAR-10.

Things you should try:

- Filter size: Above we used 7x7; but VGGNet and onwards showed stacks of 3x3 filters are good.
- Number of filters: Above we used 32 filters. Do more or fewer do better?
- Batch normalization: Try adding spatial batch normalization after convolution layers and vanilla batch normalization after affine layers. Do your networks train faster?
- Network architecture: Can a deeper CNN do better? Consider these architectures:
 - [conv-relu-pool]xN - conv - relu - [affine]xM - [softmax or SVM]
 - [conv-relu-pool]xN - [affine]xM - [softmax or SVM]
 - [conv-relu-conv-relu-pool]xN - [affine]xM - [softmax or SVM]

Tips for training

For each network architecture that you try, you should tune the learning rate and regularization strength. When doing this there are a couple important things to keep in mind:

- If the parameters are working well, you should see improvement within a few hundred iterations
- Remember the coarse-to-fine approach for hyperparameter tuning: start by testing a large range of hyperparameters for just a few training iterations to find the combinations of parameters that are working at all.
- Once you have found some sets of parameters that seem to work, search more finely around these parameters. You may need to train for more epochs.

```
In [33]: # ===== #
# YOUR CODE HERE:
#   Implement a CNN to achieve greater than 65% validation accuracy
#   on CIFAR-10.
# ===== #
model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001,
    filter_size = 7, num_filters = 64)

solver = Solver(model, data,
    num_epochs=5, batch_size=500,
    update_rule='adam',
    optim_config={
        'learning_rate': 1e-3,
    },
    verbose=True, print_every=10)
solver.train()

# ===== #
# END YOUR CODE HERE
# ===== #
```

```
(Iteration 1 / 490) loss: 2.306706
(Epoch 0 / 5) train acc: 0.093000; val_acc: 0.112000
(Iteration 11 / 490) loss: 2.087431
(Iteration 21 / 490) loss: 1.929070
(Iteration 31 / 490) loss: 1.751228
(Iteration 41 / 490) loss: 1.719299
(Iteration 51 / 490) loss: 1.587831
(Iteration 61 / 490) loss: 1.502691
(Iteration 71 / 490) loss: 1.549977
(Iteration 81 / 490) loss: 1.402410
(Iteration 91 / 490) loss: 1.510931
(Epoch 1 / 5) train acc: 0.511000; val_acc: 0.512000
(Iteration 101 / 490) loss: 1.440094
(Iteration 111 / 490) loss: 1.445302
(Iteration 121 / 490) loss: 1.402213
(Iteration 131 / 490) loss: 1.372276
(Iteration 141 / 490) loss: 1.343683
(Iteration 151 / 490) loss: 1.268084
(Iteration 161 / 490) loss: 1.358481
(Iteration 171 / 490) loss: 1.306516
(Iteration 181 / 490) loss: 1.231888
(Iteration 191 / 490) loss: 1.309796
(Epoch 2 / 5) train acc: 0.596000; val_acc: 0.582000
(Iteration 201 / 490) loss: 1.185613
(Iteration 211 / 490) loss: 1.254201
(Iteration 221 / 490) loss: 1.253186
(Iteration 231 / 490) loss: 1.259793
(Iteration 241 / 490) loss: 1.220255
(Iteration 251 / 490) loss: 1.127548
(Iteration 261 / 490) loss: 1.188627
(Iteration 271 / 490) loss: 1.142228
(Iteration 281 / 490) loss: 1.292445
(Iteration 291 / 490) loss: 1.178847
(Epoch 3 / 5) train acc: 0.591000; val_acc: 0.594000
(Iteration 301 / 490) loss: 1.191051
(Iteration 311 / 490) loss: 1.059465
(Iteration 321 / 490) loss: 0.995821
(Iteration 331 / 490) loss: 1.034540
(Iteration 341 / 490) loss: 0.979557
(Iteration 351 / 490) loss: 1.033804
(Iteration 361 / 490) loss: 1.093286
(Iteration 371 / 490) loss: 0.934026
(Iteration 381 / 490) loss: 0.986693
(Iteration 391 / 490) loss: 1.040440
(Epoch 4 / 5) train acc: 0.682000; val_acc: 0.616000
(Iteration 401 / 490) loss: 1.050200
(Iteration 411 / 490) loss: 0.997861
(Iteration 421 / 490) loss: 0.988884
(Iteration 431 / 490) loss: 0.912053
(Iteration 441 / 490) loss: 1.006697
(Iteration 451 / 490) loss: 0.964885
(Iteration 461 / 490) loss: 0.988789
(Iteration 471 / 490) loss: 0.996586
(Iteration 481 / 490) loss: 0.931486
(Epoch 5 / 5) train acc: 0.696000; val_acc: 0.638000
```

```
In [34]: model = ThreeLayerConvNet(weight_scale=0.001, hidden_dim=500, reg=0.001,
    filter_size = 7, num_filters = 64)

solver = Solver(model, data,
                num_epochs=7, batch_size=500,
                update_rule='adam',
                optim_config={
                    'learning_rate': 5e-4,
                },
                verbose=True, print_every=10)
solver.train()
```



```
(Iteration 1 / 686) loss: 2.306666
(Epoch 0 / 7) train acc: 0.123000; val_acc: 0.130000
(Iteration 11 / 686) loss: 1.948879
(Iteration 21 / 686) loss: 1.809694
(Iteration 31 / 686) loss: 1.653385
(Iteration 41 / 686) loss: 1.501529
(Iteration 51 / 686) loss: 1.461271
(Iteration 61 / 686) loss: 1.359689
(Iteration 71 / 686) loss: 1.391672
(Iteration 81 / 686) loss: 1.474507
(Iteration 91 / 686) loss: 1.489841
(Epoch 1 / 7) train acc: 0.512000; val_acc: 0.539000
(Iteration 101 / 686) loss: 1.423623
(Iteration 111 / 686) loss: 1.258785
(Iteration 121 / 686) loss: 1.326817
(Iteration 131 / 686) loss: 1.263780
(Iteration 141 / 686) loss: 1.300442
(Iteration 151 / 686) loss: 1.311888
(Iteration 161 / 686) loss: 1.197304
(Iteration 171 / 686) loss: 1.210910
(Iteration 181 / 686) loss: 1.188379
(Iteration 191 / 686) loss: 1.084884
(Epoch 2 / 7) train acc: 0.606000; val_acc: 0.609000
(Iteration 201 / 686) loss: 1.143054
(Iteration 211 / 686) loss: 1.212496
(Iteration 221 / 686) loss: 1.189139
(Iteration 231 / 686) loss: 1.031809
(Iteration 241 / 686) loss: 1.091389
(Iteration 251 / 686) loss: 1.053163
(Iteration 261 / 686) loss: 1.130875
(Iteration 271 / 686) loss: 1.083643
(Iteration 281 / 686) loss: 1.082385
(Iteration 291 / 686) loss: 1.063271
(Epoch 3 / 7) train acc: 0.667000; val_acc: 0.610000
(Iteration 301 / 686) loss: 1.096024
(Iteration 311 / 686) loss: 1.017514
(Iteration 321 / 686) loss: 0.983623
(Iteration 331 / 686) loss: 0.862629
(Iteration 341 / 686) loss: 1.040473
(Iteration 351 / 686) loss: 1.014821
(Iteration 361 / 686) loss: 0.950051
(Iteration 371 / 686) loss: 0.857804
(Iteration 381 / 686) loss: 0.932675
(Iteration 391 / 686) loss: 0.904028
(Epoch 4 / 7) train acc: 0.671000; val_acc: 0.622000
(Iteration 401 / 686) loss: 0.916411
(Iteration 411 / 686) loss: 1.070502
(Iteration 421 / 686) loss: 0.919669
(Iteration 431 / 686) loss: 0.891754
(Iteration 441 / 686) loss: 0.859301
(Iteration 451 / 686) loss: 0.809891
(Iteration 461 / 686) loss: 0.794971
(Iteration 471 / 686) loss: 0.884732
(Iteration 481 / 686) loss: 0.863436
(Epoch 5 / 7) train acc: 0.735000; val_acc: 0.636000
(Iteration 491 / 686) loss: 0.867697
(Iteration 501 / 686) loss: 0.811700
```

```
(Iteration 511 / 686) loss: 0.785466
(Iteration 521 / 686) loss: 0.864686
(Iteration 531 / 686) loss: 0.709413
(Iteration 541 / 686) loss: 0.813629
(Iteration 551 / 686) loss: 0.778565
(Iteration 561 / 686) loss: 0.746269
(Iteration 571 / 686) loss: 0.772990
(Iteration 581 / 686) loss: 0.644451
(Epoch 6 / 7) train acc: 0.767000; val_acc: 0.657000
(Iteration 591 / 686) loss: 0.825515
(Iteration 601 / 686) loss: 0.680056
(Iteration 611 / 686) loss: 0.702553
(Iteration 621 / 686) loss: 0.750595
(Iteration 631 / 686) loss: 0.738825
(Iteration 641 / 686) loss: 0.751157
(Iteration 651 / 686) loss: 0.703667
(Iteration 661 / 686) loss: 0.809817
(Iteration 671 / 686) loss: 0.734236
(Iteration 681 / 686) loss: 0.670993
(Epoch 7 / 7) train acc: 0.803000; val_acc: 0.663000
```

Therefore, validation accuracy of > 65% achieved.

```

1 import numpy as np
2 from nndl.layers import *
3 import pdb
4
5 """
6 This code was originally written for CS 231n at Stanford University
7 (cs231n.stanford.edu). It has been modified in various areas for use in the
8 ECE 239AS class at UCLA. This includes the descriptions of what code to
9 implement as well as some slight potential changes in variable names to be
10 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
11 permission to use this code. To see the original version, please visit
12 cs231n.stanford.edu.
13 """
14
15 def conv_forward_naive(x, w, b, conv_param):
16     """
17     A naive implementation of the forward pass for a convolutional layer.
18
19     The input consists of N data points, each with C channels, height H and width
20     W. We convolve each input with F different filters, where each filter spans
21     all C channels and has height HH and width WW.
22
23     Input:
24     - x: Input data of shape (N, C, H, W)
25     - w: Filter weights of shape (F, C, HH, WW)
26     - b: Biases, of shape (F,)
27     - conv_param: A dictionary with the following keys:
28       - 'stride': The number of pixels between adjacent receptive fields in the
29         horizontal and vertical directions.
30       - 'pad': The number of pixels that will be used to zero-pad the input.
31
32     Returns a tuple of:
33     - out: Output data, of shape (N, F, H', W') where H' and W' are given by
34       H' = 1 + (H + 2 * pad - HH) / stride
35       W' = 1 + (W + 2 * pad - WW) / stride
36     - cache: (x, w, b, conv_param)
37     """
38
39     pad = conv_param['pad']
40     stride = conv_param['stride']
41
42     # ===== #
43     # YOUR CODE HERE:
44     # Implement the forward pass of a convolutional neural network.
45     # Store the output as 'out'.
46     # Hint: to pad the array, you can use the function np.pad.
47     # ===== #
48     (N, C, H, W) = x.shape
49     (F, C, HH, WW) = w.shape
50     H_new = 1 + (H + 2 * pad - HH) / stride
51     W_new = 1 + (W + 2 * pad - WW) / stride
52     xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
53     out = np.zeros((N,F,H_new, W_new))
54
55     for n in range(N):
56         for f in range(F):
57             for h in range(H_new):
58                 for wd in range(W_new):
59                     h1 = h*stride
60                     h2 = h1+HH
61                     w1 = wd*stride
62                     w2 = w1+WW
63                     window = xpad[n,:,h1:h2, w1:w2] * w[f,:,:,:]
64                     sumWindow = np.sum(window)
65                     out[n,f,h,wd] = sumWindow + b[f]
66
67
68
69     # ===== #
70     # END YOUR CODE HERE
71     # ===== #
72
73     cache = (x, w, b, conv_param)
74     return out, cache

```

```

75
76
77 def conv_backward_naive(dout, cache):
78     """
79     A naive implementation of the backward pass for a convolutional layer.
80
81     Inputs:
82     - dout: Upstream derivatives.
83     - cache: A tuple of (x, w, b, conv_param) as in conv_forward_naive
84
85     Returns a tuple of:
86     - dx: Gradient with respect to x
87     - dw: Gradient with respect to w
88     - db: Gradient with respect to b
89     """
90     N, F, out_height, out_width = dout.shape
91     x, w, b, conv_param = cache
92     dx, dw, db = None, None, None
93     dx = np.zeros_like(x)
94     dw = np.zeros_like(w)
95     db = np.zeros_like(b)
96
97     N, C, H, W = x.shape
98     F, _, HH, WW = w.shape
99
100    stride, pad = [conv_param['stride'], conv_param['pad']]
101    xpad = np.pad(x, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
102    dxpad = np.pad(dx, ((0,0), (0,0), (pad,pad), (pad,pad)), mode='constant')
103    num_filts, _, f_height, f_width = w.shape
104
105    for n in range(N):
106        for f in range(F):
107            for ht in range(out_height):
108                for wd in range(out_width):
109                    h1 = ht*stride
110                    h2 = h1+HH
111                    w1 = wd*stride
112                    w2 = w1+WW
113                    dxpad[n,:,h1:h2, w1:w2] += w[f, :, :, :] * dout[n, f, ht, wd]
114                    dw[f, :, :, :] += xpad[n, :, h1:h2, w1:w2] * dout[n, f, ht, wd]
115                    db[f] += dout[n, f, ht, wd]
116
117    dx[n, :, :, :] = dxpad[n, :, pad:-pad, pad:-pad]
118
119
120    # ===== #
121    # YOUR CODE HERE:
122    # Implement the backward pass of a convolutional neural network.
123    # Calculate the gradients: dx, dw, and db.
124    # ===== #
125
126
127    # ===== #
128    # END YOUR CODE HERE
129    # ===== #
130
131    return dx, dw, db
132
133
134 def max_pool_forward_naive(x, pool_param):
135     """
136     A naive implementation of the forward pass for a max pooling layer.
137
138     Inputs:
139     - x: Input data, of shape (N, C, H, W)
140     - pool_param: dictionary with the following keys:
141         - 'pool_height': The height of each pooling region
142         - 'pool_width': The width of each pooling region
143         - 'stride': The distance between adjacent pooling regions
144
145     Returns a tuple of:
146     - out: Output data
147     - cache: (x, pool_param)
148     """
149     out = None

```

```

150 N, C, H, W = x.shape
151 pool_height = pool_param['pool_height']
152 pool_width = pool_param['pool_width']
153 stride = pool_param['stride']
154 H_new = 1 + (H - pool_height) / stride
155 W_new = 1 + (W - pool_width) / stride
156 out = np.zeros((N,C,H_new, W_new))
157 # ===== #
158 # YOUR CODE HERE:
159 # Implement the max pooling forward pass.
160 # ===== #
161 for n in range(N):
162     for c in range(C):
163         for h in range(H_new):
164             for wd in range(W_new):
165                 h1 = h*stride
166                 h2 = h1+pool_height
167                 w1 = wd*stride
168                 w2 = w1+pool_width
169                 window = np.max(x[n,c,h1:h2, w1:w2])
170                 out[n,c,h,wd] = window
171
172 # ===== #
173 # END YOUR CODE HERE
174 # ===== #
175 cache = (x, pool_param)
176 return out, cache
177
178 def max_pool_backward_naive(dout, cache):
179     """
180     A naive implementation of the backward pass for a max pooling layer.
181
182     Inputs:
183     - dout: Upstream derivatives
184     - cache: A tuple of (x, pool_param) as in the forward pass.
185
186     Returns:
187     - dx: Gradient with respect to x
188     """
189     dx = None
190     x, pool_param = cache
191     pool_height, pool_width, stride = pool_param['pool_height'], pool_param['pool_width'], pool_param['stride']
192
193     # ===== #
194     # YOUR CODE HERE:
195     # Implement the max pooling backward pass.
196     # ===== #
197     N, F, out_height, out_width = dout.shape
198     dx = np.zeros_like(x)
199
200     N, C, H, W = x.shape
201
202
203     for n in range(N):
204         for c in range(C):
205             for ht in range(out_height):
206                 for wd in range(out_width):
207                     h1 = ht*stride
208                     h2 = h1+pool_height
209                     w1 = wd*stride
210                     w2 = w1+pool_width
211                     window = x[n,c,h1:h2, w1:w2]
212                     window2 = np.reshape(window, (pool_height*pool_width))
213                     window3 = np.zeros_like(window2)
214                     window3[np.argmax(window2)] = 1
215
216                     dx[n,c,h1:h2, w1:w2] = np.reshape(window3, (pool_height, pool_width)) * dout[n,c,ht,wd]
217
218
219
220 # ===== #
221 # END YOUR CODE HERE
222 # ===== #
223
224 return dx

```

```

225
226 def spatial_batchnorm_forward(x, gamma, beta, bn_param):
227     """
228     Computes the forward pass for spatial batch normalization.
229
230     Inputs:
231     - x: Input data of shape (N, C, H, W)
232     - gamma: Scale parameter, of shape (C,)
233     - beta: Shift parameter, of shape (C,)
234     - bn_param: Dictionary with the following keys:
235       - mode: 'train' or 'test'; required
236       - eps: Constant for numeric stability
237       - momentum: Constant for running mean / variance. momentum=0 means that
238         old information is discarded completely at every time step, while
239         momentum=1 means that new information is never incorporated. The
240         default of momentum=0.9 should work well in most situations.
241       - running_mean: Array of shape (D,) giving running mean of features
242       - running_var: Array of shape (D,) giving running variance of features
243
244     Returns a tuple of:
245     - out: Output data, of shape (N, C, H, W)
246     - cache: Values needed for the backward pass
247     """
248     out, cache = None, None
249     N, C, H, W = x.shape
250     XTranspose = x.transpose(0,2,3,1)
251     x_reshape = np.reshape(XTranspose, (N*H*W, C))
252     # ===== #
253     # YOUR CODE HERE:
254     # Implement the spatial batchnorm forward pass.
255     #
256     # You may find it useful to use the batchnorm forward pass you
257     # implemented in HW #4.
258     # ===== #
259     out1, cache = batchnorm_forward(x_reshape, gamma, beta, bn_param)
260
261     out = out1.reshape((N,H,W,C)).transpose(0,3,1,2)
262
263     # ===== #
264     # END YOUR CODE HERE
265     # ===== #
266
267     return out, cache
268
269
270 def spatial_batchnorm_backward(dout, cache):
271     """
272     Computes the backward pass for spatial batch normalization.
273
274     Inputs:
275     - dout: Upstream derivatives, of shape (N, C, H, W)
276     - cache: Values from the forward pass
277
278     Returns a tuple of:
279     - dx: Gradient with respect to inputs, of shape (N, C, H, W)
280     - dgamma: Gradient with respect to scale parameter, of shape (C,)
281     - dbeta: Gradient with respect to shift parameter, of shape (C,)
282     """
283     dx, dgamma, dbeta = None, None, None
284
285     # ===== #
286     # YOUR CODE HERE:
287     # Implement the spatial batchnorm backward pass.
288     #
289     # You may find it useful to use the batchnorm forward pass you
290     # implemented in HW #4.
291     # ===== #
292     dx = np.zeros_like(dout)
293     N, C, H, W = dout.shape
294     doutTranspose = dout.transpose((0, 2, 3, 1))
295     dout_reshape = np.reshape(doutTranspose, (-1, C))
296     dx_flat, dgamma, dbeta = batchnorm_backward(dout_reshape, cache)
297     dx = dx_flat.reshape((N, H, W, C)).transpose(0, 3, 1, 2)
298
299     # ===== #

```

```
300  # END YOUR CODE HERE
301  # ===== #
302
303  return dx, dgamma, dbeta
```

```

1 import numpy as np
2
3 from nndl.layers import *
4 from nndl.conv_layers import *
5 from cs231n.fast_layers import *
6 from nndl.layer_utils import *
7 from nndl.conv_layer_utils import *
8
9 import pdb
10
11 """
12 This code was originally written for CS 231n at Stanford University
13 (cs231n.stanford.edu). It has been modified in various areas for use in the
14 ECE 239AS class at UCLA. This includes the descriptions of what code to
15 implement as well as some slight potential changes in variable names to be
16 consistent with class nomenclature. We thank Justin Johnson & Serena Yeung for
17 permission to use this code. To see the original version, please visit
18 cs231n.stanford.edu.
19 """
20
21 class ThreeLayerConvNet(object):
22     """
23     A three-layer convolutional network with the following architecture:
24
25     conv - relu - 2x2 max pool - affine - relu - affine - softmax
26
27     The network operates on minibatches of data that have shape (N, C, H, W)
28     consisting of N images, each with height H and width W and with C input
29     channels.
30     """
31
32     def __init__(self, input_dim=(3, 32, 32), num_filters=32, filter_size=7,
33                 hidden_dim=100, num_classes=10, weight_scale=1e-3, reg=0.0,
34                 dtype=np.float32, use_batchnorm=False):
35         """
36         Initialize a new network.
37
38         Inputs:
39         - input_dim: Tuple (C, H, W) giving size of input data
40         - num_filters: Number of filters to use in the convolutional layer
41         - filter_size: Size of filters to use in the convolutional layer
42         - hidden_dim: Number of units to use in the fully-connected hidden layer
43         - num_classes: Number of scores to produce from the final affine layer.
44         - weight_scale: Scalar giving standard deviation for random initialization
45           of weights.
46         - reg: Scalar giving L2 regularization strength
47         - dtype: numpy datatype to use for computation.
48         """
49         self.use_batchnorm = use_batchnorm
50         self.params = {}
51         self.reg = reg
52         self.dtype = dtype
53
54         # ===== #
55         # YOUR CODE HERE:
56         # Initialize the weights and biases of a three layer CNN. To initialize:
57         # - the biases should be initialized to zeros.
58         # - the weights should be initialized to a matrix with entries
59         #   drawn from a Gaussian distribution with zero mean and
60         #   standard deviation given by weight_scale.
61         # ===== #
62
63         C, H, W = input_dim
64         F = num_filters
65         filterHeight = filter_size
66         filterWidth = filter_size
67         stride = 1
68         P = (filter_size - 1) / 2
69         Hc = ((H + 2 * P - filterHeight) / stride) + 1

```



```

70 Wc = ((W + 2 * P - filterWidth) / stride) + 1
71
72 W1 = weight_scale * np.random.randn(F, C, filterHeight, filterWidth)
73 b1 = np.zeros((F))
74
75
76 width_pool = 2
77 height_pool = 2
78 # stride_pool = 2
79 # Hp = ((Hc - height_pool) / stride_pool) + 1
80 # Wp = ((Wc - width_pool) / stride_pool) + 1
81
82
83
84 Hh = hidden_dim
85 W2 = weight_scale * np.random.randn((F * Hc * Wc)/(width_pool* height_pool), Hh)
86 b2 = np.zeros((Hh))
87
88
89
90 Hc = num_classes
91 W3 = weight_scale * np.random.randn(Hh, Hc)
92 b3 = np.zeros((Hc))
93
94 self.params['W1'], self.params['b1'] = W1, b1
95 self.params['W2'], self.params['b2'] = W2, b2
96 self.params['W3'], self.params['b3'] = W3, b3
97 # ===== #
98 # END YOUR CODE HERE
99 # ===== #
100
101 for k, v in self.params.items():
102     self.params[k] = v.astype(dtype)
103
104
105 def loss(self, X, y=None):
106     """
107     Evaluate loss and gradient for the three-layer convolutional network.
108
109     Input / output: Same API as TwoLayerNet in fc_net.py.
110     """
111     W1, b1 = self.params['W1'], self.params['b1']
112     W2, b2 = self.params['W2'], self.params['b2']
113     W3, b3 = self.params['W3'], self.params['b3']
114
115     # pass conv_param to the forward pass for the convolutional layer
116     filter_size = W1.shape[2]
117     conv_param = {'stride': 1, 'pad': (filter_size - 1) / 2}
118
119     # pass pool_param to the forward pass for the max-pooling layer
120     pool_param = {'pool_height': 2, 'pool_width': 2, 'stride': 2}
121
122     scores = None
123
124     # ===== #
125     # YOUR CODE HERE:
126     # Implement the forward pass of the three layer CNN. Store the output
127     # scores as the variable "scores".
128     # ===== #
129     out1, cache1 = conv_relu_pool_forward(X, W1, b1, conv_param, pool_param)
130     out2, cache2 = affine_relu_forward(out1, W2, b2)
131     scores, cache3 = affine_forward(out2, W3, b3)
132
133
134     # N, F, Hp, Wp = out1.shape
135     # out2 = out1.reshape((N, F * Hp * Wp))
136
137
138     # ===== #
139     # END YOUR CODE HERE

```

```

140 # ===== #
141
142 if y is None:
143     return scores
144
145 loss, grads = 0, {}
146 # ===== #
147 # YOUR CODE HERE:
148 # Implement the backward pass of the three layer CNN. Store the grads
149 # in the grads dictionary, exactly as before (i.e., the gradient of
150 # self.params[k] will be grads[k]). Store the loss as "loss", and
151 # don't forget to add regularization on ALL weight matrices.
152 # ===== #
153 loss, dscores = softmax_loss(scores, y)
154 reg_loss = 0.5 * self.reg * np.sum(W1**2)
155 reg_loss += 0.5 * self.reg * np.sum(W2**2)
156 reg_loss += 0.5 * self.reg * np.sum(W3**2)
157 loss = loss + reg_loss
158
159 dx3, grads['W3'], grads['b3'] = affine_backward(dscores, cache3)
160 dx2, grads['W2'], grads['b2'] = affine_relu_backward(dx3, cache2)
161 dx1, grads['W1'], grads['b1'] = conv_relu_pool_backward(dx2, cache1)
162
163 grads['W1'] += self.reg * self.params['W1']
164 grads['W2'] += self.reg * self.params['W2']
165 grads['W3'] += self.reg * self.params['W3']
166 # ===== #
167 # END YOUR CODE HERE
168 # ===== #
169
170 return loss, grads
171
172
173 pass

```